

STRATEGIC DEAL INTELLIGENCE ASSISTANT

Agentic AI Engineering Walkthrough

Grounded • Permission-Aware • Observable • Human-Approved




Oded Goldberg

Candidate — Agentic AI Engineer

FROM FRAGMENTED CONTEXT TO DECISION SUPPORT

The challenge is not missing data. It is disconnected context.



Faster preparation. Better grounding. Safer decisions.

LOCAL PROTOTYPE ARCHITECTURE

A governed agent workflow around live LLM reasoning

CORE PRINCIPLES


- 01

Authorization
Before Retrieval
- 02

Specialized
Agents
- 03

Human Approval
for Sensitive Actions
- 04

Observable by Design
Metrics • Traces • Audit Trail



AUTHORIZATION BEFORE EVIDENCE REACHES THE MODEL

Security is enforced before retrieval — not delegated to the prompt.

SECURITY PRINCIPLES



Authorization Before Retrieval



Permission Aware Filtering


User • Opportunity • Source • Metadata



Unauthorized Restricted Data Never Enters Model Context



Clean Denial • No Partial Answer



SPECIALIZED AGENTS → ONE STRATEGIC BRIEF

Focused responsibilities make reasoning easier to control, validate and evaluate.

01

Deal Context Agent

Canonical CRM facts

Opportunity • Stage • Amount • Close Date • Owner

02

Conversation Intelligence Agent

Goals • Objections • Urgency • Competitors • Action Items

03

Stakeholder Map Agent

Economic Buyer • Champion • Blockers •
Procurement • Unknowns

04


Negotiation Strategy Agent

Prioritized Next Actions • Owners • Rationale



EVALUATION — CAN I TRUST THE BEHAVIOR?

I evaluate the system around business critical failure modes, not only answer quality.



FINAL RUN RESULTS	Golden Tests	Unauthorized Leakage	Citation Validation	Required Slack Evidence
	10/10 passed	0 retricted evidence leaked	10/10 citations validated	7 synthetic updates indexed Slack citations verified

A CONVINCING 11M RESPONSE IS INSUFFICIENT — CRITICAL SYSTEM BEHAVIOR MUST BE MEASURABLE AND REGRESSION-TESTED.

WHAT BREAKS IN PRODUCTION?

The prototype validates the workflow. Production adds a durable data plane, incremental ingest, security, and operational scale.

PROTOTYPE		PRODUCTION EVOLUTION	
CURRENT / LOCAL RUNNABLE PROTOTYPE		FUTURE / PROPOSED	
	Local single-process application	Data Ingestion	Connectors, webhooks, polling, SQS/EventBridge
	Local Qdrant index	Index Lifecycle	Incremental upsets, deletes, replay, idempotency
	Local workflow state	Freshness & History	Hot index for recent data, archive for historical search
	Local run artifacts	Durable State	PostgreSQL workflow state, approvals, ingest checkpoints
	Static file ingest	Security	API authentication, IAM, secrets management; tenant isolation
	Single application instance	Availability & Scale	Retries, timeouts, rate limits, dead-letter queues
	Live LLM provider	Observability	Centralized logs, metrics, traces, alerts, multi-worker capacity planning



LIVE DEMO – THREE REQUESTS, THREE DELIBERATE OUTCOMES

01

AUTHORIZED STANDARD DEAL

USR-5001 → OPP-1001

Show: Brief Generation Retrieved Evidence Citations Recommended Next Actions



Expected Outcome:



GROUNDING BRIEF

02

AUTHORIZED RESTRICTED DEAL

USR-5003 → OPP-1003

Show: Restricted Evidence Sensitive Recommendation Approval Flow



Expected Outcome:



HUMAN REVIEW REQUIRED

03

UNAUTHORIZED REQUEST

USR-5007 → OPP-1003

Show: Authorization Failure Authorization Failure



Expected Outcome:



DENIED BEFORE RETRIEVAL



PROOF OF MULTI-SOURCE GROUNDING



Generated
Slack Update



Retrieved
Evidence



Cited Finding
in the Brief



QUESTIONS?